

# The AI Visibility Evidence Model.

Five Publisher-Side Factors, Graded by Evidence.

## Stefan Petschinka

ORCID: 0009-0009-8223-0339

Published by: richresults.ai

Version: 1.0

First published: July 30, 2026

Revised: August 30, 2026

DOI: 10.5281/zenodo.21766951

License: Creative Commons Attribution 4.0 International (CC BY 4.0)

<https://www.richresults.ai/evidence.html>

---

## RECOMMENDED CITATION

Petschinka, Stefan (2026). *The AI Visibility Evidence Model*. Version 1.0. DOI: 10.5281/zenodo.21766951.

## What the AI Visibility Evidence Model is.

The AI Visibility Evidence Model is a reference model that orders the publisher-side factors behind AI visibility by strength of evidence. It defines five factors: Topical Relevance, Machine Access, Entity Consistency, Extractability and Independent Corroboration, and assigns each a documented evidence grade based on peer-reviewed research, controlled preprints and official platform documentation.

The model exists because the field still lacks a concise publisher-side reference that maps the main actionable factors to explicit evidence grades and primary sources. Every factor in the model carries its grade and its sources, so every statement in this document can be checked against the original studies and official documentation listed below.

## A map of the evidence, not a methodology.

The AI Visibility Evidence Model maps what the evidence shows works. The AEO Mastery Framework describes how richresults.ai implements it. The model is descriptive: it reports the state of the research. The framework is prescriptive: it defines a working method. Neither replaces the other.

The model provides an evidence-based order for publisher-side work on AI visibility. It shows which factors are supported as drivers, which operate as documented prerequisites or error-reduction mechanisms, and which claims are not supported by current evidence. AI visibility remains a distribution across repeated, non-deterministic answers, so the model is not a ranking system, a score or a promise of a specific citation. Its practical value is priority: it shows where publishers can act on evidence-backed mechanisms and how strongly each mechanism is documented.

One boundary is deliberate: the model orders the publisher-side factors that can be worked on before and around retrieval. The retrieval stage itself, which engine selects which sources, how it ranks them and where it places them in the model context, is system-side. Controlled work shows that this stage strongly shapes citation outcomes [2], and platform documentation describes engine-specific retrieval decisions, including when a system grounds at all [13]. All five factors are publisher-side inputs into that process. The model therefore separates what publishers can improve from the engine's final retrieval decision without treating the absence of control over that final decision as evidence that publisher-side work is ineffective.

## Four grades, defined before use.

**Grade A:** peer-reviewed and controlled, or independently replicated.

**Grade B:** controlled with limited transferability to open production systems, or an official platform statement.

**Grade C:** correlational or triangulated across independent datasets, without causal proof.

**Grade D:** unsupported or contradicted by evidence.

Each factor additionally carries its mechanism type. A gate is a binary precondition, a driver influences outcomes gradually, and hygiene reduces errors or ambiguity. Mechanism type and evidence grade are separate dimensions: a factor can be a hard gate on weak empirical evidence, or a soft driver on strong evidence. The hygiene label does not mean that a factor creates no visibility benefit; it means that the documented mechanism in this model is primarily error reduction or disambiguation, while any independent visibility uplift has not been isolated by the cited evidence.

These four grades are this model's own scale, and it is deliberately conservative. A factor rated C can still be practically relevant: the available evidence is correlational, triangulated or transferred from benchmark settings rather than causally isolated in open production systems. Grade C therefore limits the strength of the claim; it does not turn missing causal proof into evidence of no effect. Grades move as evidence accumulates, and this document carries a visible revision date for that reason.

### Factor 1: Topical Relevance

Content that directly addresses the actual question is the strongest documented content-side driver of citation. In the largest controlled citation study to date, 252,000 trials across six language models, topic match to the query and position in the model context were the dominant factors, and off-topic content was practically never cited first [1]. Controlled work confirms the same pattern from the model side: when weighing conflicting evidence, models rely heavily on a page's relevance to the query while largely ignoring stylistic authority signals such as scientific-looking references or neutral tone [14]. No entity work, no markup and no authority signal compensates for content that does not answer the question being asked.

**Mechanism: driver. Evidence grade: A.** Peer-reviewed, controlled, convergent across models and study designs [1, 2, 14].

## Factor 2: Machine Access

A source that crawlers cannot reach cannot be retrieved, and a source that cannot be retrieved cannot be used for the content of an answer. Machine Access covers crawl permissions for the relevant bots, index presence, crawlable and renderable main content, and firewall configurations that do not silently block AI crawlers. Platform documentation is explicit on both sides of this gate. OpenAI requires OAI-SearchBot access for a site's content to be used in ChatGPT search answers; excluded pages can still appear as navigational links [12]. Google requires indexed, snippet-eligible pages and states that its AI features run on the same index and ranking systems as classical search [11].

**Mechanism: gate. Evidence grade: B, official platform documentation.** Without access, a page's content cannot be retrieved for use in an answer; Machine Access is therefore a documented prerequisite [11, 12]. The cited sources do not establish that access alone guarantees retrieval or citation, and they do not quantify an independent visibility uplift from access itself.

## Factor 3: Entity Consistency

Consistent entity signals support attributing relevant content to the right source. Structured data, stable identifiers, canonical name strings and connected external profiles can reduce ambiguity in how systems resolve who is speaking. Google states that no special markup is required for its AI features [11]; this establishes that special markup is not a participation requirement, not that structured entity signals have no effect. Controlled work shows that knowledge-graph grounding reduces entity disambiguation errors in benchmark settings [15]. The evidence therefore supports Entity Consistency as a mechanism for clearer entity resolution and attribution. A direct causal effect on citation or mention rates in production answer engines has not been isolated, so this model does not assign an independent citation uplift to the factor.

**Mechanism: hygiene. Evidence grade: C.** Supported for disambiguation and entity resolution [11, 15]; no isolated causal citation or mention uplift in production answer engines.

## Factor 4: Extractability

A model can only cite what it can extract. The evidence supports several concrete properties of extractable information: the position of information in the model context changes outcomes causally [3, 4], and pages containing concrete numbers, definitions, comparisons and procedures show substantially higher influence on generated answers than pages without them [7]. This supports Extractability as a practical publisher-side objective: make relevant facts explicit, specific and easy to isolate once a source has been retrieved. The boundary is that these findings do not establish a universal page-formatting formula. The influence finding in production is descriptive and correlational [7], question-and-answer formatting alone does not help [7], and content rewriting tricks show no reliable effect and are frequently harmful under competition [2]. A publisher also does not control which passage a retriever selects or where that chunk lands in the model context, so answer-first page structure remains a reasoned publisher tactic rather than a demonstrated position lever.

**Mechanism: driver. Evidence grades, split:** B for causal position effects and evidence density [3, 4, 6]; C for publisher-side transfer into page structure and production environments [7], bounded by [2, 9].

## Factor 5: Independent Corroboration

Mentions beyond a publisher's own properties are supported by two relevant evidence paths. First, controlled research shows that the frequency of entity-related evidence across training documents causally affects what a model knows about the entity without searching [5]. This establishes repeated external evidence as a meaningful input into the parametric layer, although the study does not isolate the independence of those documents. Second, a preprint analysis reports that several AI search systems show a strong preference for earned media over brand-owned content, based on the authors' own source classification [8]. Together, these findings support external corroboration as a practical publisher-side objective: important entity claims are stronger when they are repeated and supported beyond the publisher's own site. The narrower claim that independence or authenticity by itself causes higher AI visibility has not been isolated by either source, so this model does not assign a separate causal uplift to independence alone. Manufactured or self-produced mentions remain unsupported as a positive visibility tactic.

**Mechanism: driver for the parametric layer. Evidence grades, split:** B for training-data frequency [5]; C for observed earned-media preference [8]. Independence itself has not been isolated as a separate causal visibility factor.

## Specific claims graded D, for different reasons.

**llms.txt as a visibility lever.** Google explicitly states that it does not use llms.txt for search or for generative search features [11]. Within the source register reviewed here, no platform documents the file as a ranking or visibility signal.

**Content rewriting as a reliable citation lever.** The broadest controlled benchmark found most tested conversational optimization methods ineffective and frequently harmful to citation ranking, while classical retrieval position dominated [2]. The Grade D assessment applies to the claim that rewriting alone provides a reliable citation advantage, not to topical relevance or extractable, evidence-rich content, which are evaluated separately in this model.

**Schema as a direct citation switch.** The evidence does not support treating structured data as a guaranteed or isolated citation trigger. Google states that special markup is not required for its AI features [11], and the sources reviewed here do not establish a causal citation uplift from schema alone. This does not contradict the role of structured data within Entity Consistency: structured data can clarify and disambiguate entity signals, as described in Factor 3. The Grade D assessment applies to the direct citation-switch claim, not to structured data as an entity-clarification mechanism.

**Single-run AI ranking positions as a stable metric.** Answers vary widely across otherwise identical runs, so a position from a single run is not a reliable rank [9]. Positions become meaningful as distributions across repeated, paraphrased measurements with uncertainty intervals; visibility is a share, not a single rank.

**Manufactured mentions.** No reliable evidence exists for a positive effect of manufactured or self-produced mentions. In the literature, the line between optimization and manipulation is not determined by effectiveness, but by truthfulness, verifiable evidence, the separation of content from model instructions, and disclosure of commercial intent [9]. What this model states about corroboration rests on training-data frequency [5] and an observed earned-media preference [8]; neither of those works examines self-produced mentions.

## A strong measurement program tracks both.

An independent audit of eight AI search engines found that the engines collectively provided incorrect answers to more than 60 percent of source-attribution queries, with premium systems frequently confidently wrong [10]. This makes citation fidelity a second measurement target alongside visibility. Visibility measures whether and how often an entity appears; fidelity measures whether the system represents that entity accurately and attributes information to the correct source. Tracking both separates successful retrieval and citation from correct representation and gives organizations a more complete picture of their AI visibility.

## How this model was compiled.

Every claim in this model was verified against its source: the papers, official documentation and datasets listed in the register below. The synthesis was additionally checked for completeness and counter-arguments by prompting several AI systems with the same evidence question without access to this model. Convergence across systems is an editorial plausibility check, not independent scientific validation: systems share training data, sources and failure modes, and can converge on the same popular error.

**Disclosure:** This reference model is published by richresults.ai and authored by Stefan Petschinka, who also developed the linked AEO Mastery Framework. The source assessment, synthesis and evidence grades are the author's own work and have not undergone external peer review.

Three conditions define how the findings in this model should be interpreted. Production systems are non-deterministic, so individual observations are measurements rather than isolated causal proof. Models and retrieval methods change, so findings carry dates and evidence grades can be updated as new evidence accumulates. Live answer engines do not currently allow a single publisher-side intervention to be causally isolated end to end; the evidence grades therefore distinguish what is directly demonstrated, what is transferred from controlled settings and what remains open. This keeps the model actionable without overstating certainty.

### ABOUT THE AUTHOR

## Stefan Petschinka.

**AEO Strategist · Entity Architect · Founder of richresults.ai**

Stefan Petschinka is an AEO Strategist, Entity Architect and founder of richresults.ai. Stefan Petschinka specializes in AI Visibility, Answer Engine Optimization, Generative Engine Optimization, Large Language Model Optimization and Entity Building. His work builds AI Visibility through Entity Building: the machine-readable identity layer that makes organizations, brands and experts understandable, citable and recommendable by ChatGPT, Perplexity, Claude, Gemini and Google AI Search.

Stefan Petschinka is the author of two nonfiction books on how AI language models answer: FEED THE MACHINE and ECHO: The Dark Psychology of AI.

Web: <https://www.richresults.ai/stefan-petschinka.html> · ORCID: 0009-0009-8223-0339

## Numbered as cited above.

1. Vishwakarma, Kumar, Jamidar (2026). What Gets Cited: Competitive GEO in AI Answer Engines. SIGIR 2026. DOI:10.1145/3805712.3808445, arXiv:2605.25517. Peer-reviewed; authors affiliated with Sprinklr.
2. Puerto et al. (2025). C-SEO Bench: Does Conversational SEO Work? NeurIPS 2025 Datasets and Benchmarks. arXiv:2506.11097.
3. Liu et al. (2024). Lost in the Middle: How Language Models Use Long Contexts. TACL 2024. arXiv:2307.03172.
4. Hsieh et al. (2024). Found in the Middle: Calibrating Positional Attention Bias. ACL 2024 Findings. arXiv:2406.16008.
5. Kandpal et al. (2023). Large Language Models Struggle to Learn Long-Tail Knowledge. ICML 2023. arXiv:2211.08411.
6. Aggarwal et al. (2024). GEO: Generative Engine Optimization. KDD 2024. arXiv:2311.09735. Effects conditional on fixed-context settings; see [2] and [9].
7. Zhang, He, Yao (2026). From Citation Selection to Citation Absorption. Preprint. arXiv:2604.25707.
8. Chen, Wang, Chen, Koudas (2025). Generative Engine Optimization: How to Dominate AI Search. Preprint, University of Toronto. arXiv:2509.08919. Acknowledges support by ktai.ai.
9. Martinez (2026). Optimizing Visibility in Generative Engines: A Critical Survey. Preprint. arXiv:2607.14035.
10. Jaźwińska, Chandrasekar (2025). AI Search Has a Citation Problem. Tow Center, Columbia Journalism Review. cjr.org.
11. Google Search Central: AI features and your website. developers.google.com. Also: Optimizing your website for generative AI features on Google Search. developers.google.com.
12. OpenAI: OAI-SearchBot documentation. openai.com/searchbot. Also: Publishers and Developers FAQ. help.openai.com.
13. Google AI for Developers: Grounding with Google Search. ai.google.dev. Documents engine-side retrieval decisions; not a publisher participation requirement.
14. Wan, Wallace, Klein (2024). What Evidence Do Language Models Find Convincing? ACL 2024, pages 7468–7484. DOI:10.18653/v1/2024.acl-long.403, arXiv:2402.11782.
15. Pons, Bilalli, Queralt (2024). Knowledge Graphs for Enhancing Large Language Models in Entity Disambiguation. Proc. 23rd Int. Semantic Web Conference (ISWC 2024), LNCS, Springer. DOI: 10.1007/978-3-031-77844-5\_9, arXiv:2505.02737.